

COWORK GENEALOGY · KICKOFF

# Junior Genealogist Assessment Kickoff

Welcome — thanks for being here. Today we get you set up; tomorrow you do the work.

# Agenda — about one hour

| Topic                                                 | Who              | ~Time    |
|-------------------------------------------------------|------------------|----------|
| The project — what we're building                     | Dallan           | 5 min    |
| The assessment — what you'll do & how you're measured | Dallan           | 10 min   |
| One-time install & the daily flow                     | Promise / Israel | ~15 min  |
| Reviewing scores in the app                           | Dallan           | ~10 min  |
| Commit & push — submitting your PR                    | Promise / Israel | ~10 min  |
| Questions                                             | Everyone         | 5–10 min |

DALLAN

5 MIN

# The Project

## What we're building

An **AI family history consultant** — a research assistant that helps people do real genealogy.

It runs inside **Claude Cowork**, a desktop app for Claude.

We make it good at genealogy by adding two things...

# Skills & Tools

## ~2 dozen Tools

Connect to outside data — search records, read a person's tree, look up places, query the FamilySearch wiki.

## ~2 dozen Skills

Do the genealogy *work* — research planning, record extraction, locality guides, conflict resolution, and more.

The skills are where the genealogical expertise lives. Making them better is the whole game.

## How we make the AI better

To improve a skill, we first have to **measure how well it works** in many different situations.

So we built a set of “**tests.**” Each test captures one realistic moment and asks: *did the AI do this well?*

A test is graded by **another AI** (the “judge”) — fast, but not reliable on genealogical judgment. That's where you come in: your corrections fix the grades *and* teach the judge to grade better over time.

## What's inside one test

- **A research scenario** — what's been found so far in this project
- **The user's input** — what the person asks the AI to do
- **Tool responses** — what the tools returned
- **The AI's response** — what the skill actually produced
- **Judge context** — instructions for how to grade the response
- **Response grades** — the judge AI's scores, which **you review**

DALLAN

10 MIN

# The Assessment



# The plan for the next several months

Working alongside our developers, we will repeatedly:

1. Run the tests for each skill
2. **Correct the grades the AI judge assigned**
3. Improve existing tests and create new ones
4. Improve the skills themselves

Correcting grades (step 2) is your core contribution — and you'll also work with the developers to improve and create tests (step 3) and improve the skills themselves (step 4).

## Where this assessment fits

- **Step 1 is already done** — the tests have been run for you.
- **Our senior genealogists have already done step 2** — they corrected the grades.
- **Your task:** redo step 2 for **4 skills** — independently correct the grades.

We then compare your corrections to the seniors' to see how well-aligned your judgment is.

Developers are here to help you get set up if you get stuck.

# How we measure success

## 1 · Aligned grades

How close are your score corrections to the **senior genealogists'** corrections?

## 2 · Low support needs

How much technical help do you need vs. working through the tasks **on your own**?

The people we hire will meet **both**: well-aligned scores *and* low support needs.

# Schedule

| When           | What                                                                                          |
|----------------|-----------------------------------------------------------------------------------------------|
| Today (Mon)    | Get all the software installed — that's what the rest of this session is for.                 |
| Tomorrow (Tue) | Correct grades on <b>4 skills</b> . We'll tell you <i>which</i> 4 first thing in the morning. |
| Tue night      | Submit a pull request with your grade corrections for those 4 skills.                         |

We're finalizing the skill list tonight — you'll get it first thing tomorrow morning.

## Time & pay

You'll be paid for **5 hours** to complete the assessment.

Today's setup + tomorrow's grading + submitting the PR all fit inside that.

# The grading vocabulary

Every test is graded on several **dimensions**:

- **Base dimensions** (always): Correctness, Completeness, Tool Arguments
- **Rubric dimensions**: 3–5 skill-specific checks

Each dimension gets one score:

**1** fail   **2** partial   **3** pass

(Tool Arguments can also be **N/A** when the skill made no tool calls.)

## Grading well

- Read the **full trace** before scoring — input, tool calls, output.
- Judge against the **rubric**, not your overall gut feeling about the skill.
- Use **2 (partial)** honestly — it isn't all-or-nothing.
- When you disagree with the AI, write **one line saying why**.
- No trick questions — grade the way your expertise tells you.

## Meet the developers

The team that built the harness and this review app — and your **technical support** throughout the assessment.

Next, the developers will walk you through the one-time install and the daily flow — live, hands-on.



DALLAN

~10 MIN

# Reviewing Scores in the App

## The review app

- A small web app that runs on your own machine.
- Launch it with `Start.bat` — a browser tab opens automatically.
- Go to `localhost:3000/results` — run logs grouped by skill.
- Click your skill's latest run to open it.

Keep the black command window open the whole time — closing it stops the app.

# Each test shows two blocks

## ① The Trace

- The skill's input
- The scenario (what's been found so far)
- Tool calls + what the tools returned
- Claude's final output

## ② The Grades

- One row per dimension
- The judge's score **1** **2** **3**
- The judge's rationale
- A picker for **your correction**

## Three things you'll do per dimension

| Situation                       | What to do                                                                                        |
|---------------------------------|---------------------------------------------------------------------------------------------------|
| You <b>agree</b> with the judge | Pick the same score; leave the comment empty.                                                     |
| You <b>disagree</b>             | Pick a different score + write a one-line reason.                                                 |
| The whole test is <b>right</b>  | Click “ <b>Agree with all</b> ” on the test header — marks every dimension reviewed in one click. |

Agreeing still counts as reviewing — you must touch *every* dimension.

# Speeding up & finishing

## Keyboard (when a picker is focused)

- `1` `2` `3` — set the score
- `Tab` / `Shift+Tab` — move between dimensions
- `?` — show shortcut help

## The completeness counter

- Header shows e.g. `8 / 12 reviewed`
- You're done when it reads **N / N**
- Your PR won't pass until **every dimension** is reviewed

Corrections save automatically — you'll see “saved” in the header.

## A fast, honest review pattern

- Scan the trace first — understand what the skill was asked and what it did.
- For each dimension, check the judge's rationale against the **rubric**.
- “Agree with all” on tests that look clearly right...
- ...and spend your attention on the ones you actually **disagree** with.

Quality over speed — but you don't have to re-justify every correct grade.

## Submitting your work

Now back to the developers — they'll show you how to **commit your corrections and create a pull request**.

That's how your graded work gets submitted for the assessment.

## If you're selected

The full three-month project is a real commitment — here's what it looks like:

- **20 hours / week**, starting **next week**.
- A **1-hour daily meeting at 12:30pm WAT**, every day **Monday–Friday**.
- You'll be assigned to a **developer team** — you'll work closely together, so you'll need to agree on **core hours when you're all online at the same time**.



5-10 MIN

# Questions?

Tomorrow morning we announce the 4 skills. Tomorrow night you submit your PR.

Thanks for participating!

We'll let you know by the end of the week whether you'll be hired for the full three-month project.